

Applied Data Science

Session 21: K-Nearest Neighbour Algorithm

Dr. Soharab Hossain Shaikh

1

K- Nearest Neighbours Algorithm (KNN)

2

Instance-Based Learning

- **Idea:**
 - Similar examples have similar label.
 - Classify new examples like similar training examples.
- **Algorithm:**
 - Given some new example x for which we need to predict its class y
 - Find most similar training examples
 - Classify x "like" these most similar examples

3


K-Nearest Neighbour Classification Algorithm

- Training method:
 - Save the training examples – No explicit training phase
- At prediction time:
 - Find the k training examples $(x_1, y_1), \dots, (x_k, y_k)$ that are closest to the test example x
 - Predict the most frequent class among those y_i 's.

4

1-Nearest Neighbour

- One of the simplest of all machine learning classifiers
- Simple idea: **label a new point** the same as the **closest known point**



6

5

1-NN – Important Aspects

A distance metric- **Euclidean**

- When different units are used for each dimension
→ Normalize each dimension by standard deviation
- For discrete data, can use hamming distance
→ $D(x1, x2)$ = number of features on which $x1$ and $x2$ differ
- Others (e.g., normal, cosine)

How many nearby neighbors to look at? **One** – for **Nearest Neighbour**

How to fit with the local points?


- Just predict the same output as the nearest neighbor.

7

6

K-Nearest Neighbour

- Generalizes 1-NN to smooth away noise in the labels
- A new point is now assigned **the most frequent label of its k nearest neighbors**



9

7


Instance-based Learning

- **Instance-based** learning is often termed *lazy* learning, as there is typically no "transformation" of training instances into more general "statements"
- Instead, the presented training data is simply stored and, when a new query instance is encountered, a set of similar, related instances is retrieved from memory and used to classify the new query instance
- Hence, **instance-based learners never form an explicit general hypothesis regarding the target function.**
- They simply compute the classification of each new query instance as needed.

10

8

Euclidean Distance – Similarity Measure



$$d = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2 + (z_1 - z_2)^2}$$

Measuring similarity with distance between the points using Euclidean method.

9

Example

- Determining **decision on scholarship** application based on the following features:
 - Household income (annual income in thousands of rupees)
 - Number of siblings in family
 - High school grade (on a scale of 1.0 – 10.0)
- Intuition (reflected on data set): Award scholarships to high-performers and to those with financial need.

11

10

Dealing with Non-Numeric Data

- Feature values are not always numbers
 - e.g. - Boolean values: Yes or No, presence or absence of an attribute
 - Categories: Colors, Gender
- Boolean values => convert to 0 or 1
 - Applies to yes-no/presence-absence attributes
- Non-binary characterizations
 - Use natural progression when applicable
 - (e.g., educational attainment: HS, College, MS, PhD => 1, 2, 3, 4, 5)
 - Assign arbitrary numbers but be careful about distances;
 - e.g., color: red, yellow, blue => 1, 2, 3

12

11

K-NN Variations

- **Value of k**
 - Larger k increases confidence in prediction
 - Note that if k is too large, decision may be skewed
- Weighted evaluation of nearest neighbors
 - Plain majority may unfairly skew decision
 - Revise algorithm so that closer neighbors have greater "vote weight"
- Other distance measures

13

12

Data Preprocessing

- Dataset may need to be preprocessed to ensure more reliable data mining results
- Conversion of non-numeric data to numeric data
- Calibration of numeric data to reduce effects of disparate ranges
- Particularly when using the Euclidean distance metric

14

13

Importance of Feature Scaling

- The Euclidean distance formula has the implicit assumption that the **different dimensions are comparable**.
- Features that span wider ranges affect the distance value more than features with limited ranges.
- Suppose household income was instead indicated in thousands of rupees per month and that grades are given on a 50-100 scale.

15

14

Standardization

- Transform raw feature values into z-scores

$$z_{ij} = \frac{x_{ij} - \mu_j}{\sigma_j}$$

x_{ij} is the value for the i^{th} sample and j^{th} feature

μ_j is the average of all x_{ij} for feature j

σ_j is the standard deviation of all x_{ij} over all input samples

- Those dimensions which have larger possible range of values will dominate the result of the distance calculation using Euclidean formula.
- To ensure all the dimensions have similar scale, we normalize the data on all the dimensions/ attributes.
- There are multiple ways of normalizing the data. We will use Z-score standardization.

16

15


Other Distance Measures

- **City-block distance (Manhattan distance)**
 - Add absolute value of differences (state-space heuristic search)
- **Cosine similarity** - Measure angle formed by the two samples (with the origin) (NLP – comparing two word-vectors)
- **Correlation-based similarity** (e.g. comparing two images)

17

16

K-NN and Irrelevant Features




- > Every feature may not be equally important.
- > Feature Selection is important – we may use weighted distance measures – give different weights to different features

Distance computation makes sense on **relevant features**.
(e.g. Decision on scholarship - high school grade and household income are better features compared to gender and geographic location of the residence of a candidate)

18

17

Importance of Feature Scaling in KNN



18

18

K-NN is Incremental

- All training instances are stored
- Model consists of the set of training instances
- Adding a new training instance only affects the computation of neighbors, which is done at execution time (i.e., lazily)

22


19

Choosing k

- Determining the optimal K is a challenge.
- Higher values of k provide smoothing - reduces the risk of overfitting due to noise in the training data
- if k is too high, we will miss out on the method's ability to capture the local structure in the data.
- In the extreme, $k = n$ = the number of records in the training dataset - all records mapped to the majority class in the training data
- In general, larger value of k suppresses impact of noise but prone to majority class dominating.
- If k is too low, we may be fitting to the noise in the data.

20

How to choose the optimum value of k?



- To get the optimal value of K, you can segregate the training and validation from the initial dataset
- Now plot the validation error (alternatively test accuracy) curve to get the optimal value of K
- This value of K should be used for all predictions

21

k-NN Time Complexity

- Suppose there are **m** instances and **n** features in the dataset
- Nearest neighbor algorithm requires computing **m** distances
- Each distance computation involves scanning through each of the **n** features
- Running **time complexity** is proportional to **$m \times n$**

19

22

Pros and Cons of KNN

Advantages

- Not impacted by Outliers (**why???**)
- Makes no assumptions about distributions of classes in feature space

Disadvantages

- Fixing the optimal value of K is a challenge.
- Does not output any models. Calculates distances for every new point.

23

Applications of KNN

- Pattern recognition in Optical character recognition
- Concept search in semantic models in NLP
- Detect similar buying patterns in customer analytics
- Text classification

24

How to *build* a KNN Classifier Model in *Scikit-Learn*

25

KNN Classifier in Scikit-Learn



`sklearn.neighbors.KNeighborsClassifier`

```
class sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, *, weights='uniform', algorithm='auto', leaf_size=30, p=2,
metric='minkowski', metric_params=None, n_jobs=None, **kwargs)
```

[\[source\]](#)

- Build an estimator model like `KNeighborsClassifier()`
- Use `fit()` function to **Train** the model with **Training Dataset**
- Use `predict()` function to **Test/Evaluate** the model with **Test Dataset**
- Use `predict()` function to make **prediction/inference** on **New Unseen Data**

26

Example Code

```
>>> X = [[0], [1], [2], [3]]
>>> y = [0, 0, 1, 1]
>>> from sklearn.neighbors import KNeighborsClassifier
>>> neigh = KNeighborsClassifier(n_neighbors=3)
>>> neigh.fit(X, y)
KNeighborsClassifier(...)
>>> print(neigh.predict([[1.1]]))
[0]
>>> print(neigh.predict_proba([[0.9]]))
[[0.66666667 0.33333333]]
```

27



28

Case-study

29

Context:

The dataset to be considered consists of a wide variety of intrusions simulated in a military network environment.

It was created in an environment to acquire raw TCP/IP dump data for a network by simulating a typical US Air Force LAN. The LAN was focused like a real environment and blasted with multiple attacks.

For each TCP/IP connection, 41 quantitative and qualitative features are obtained from normal and attack data (3 qualitative and 38 quantitative features) .

The class variable has two categories:

- Normal
- Anomalous

30

30

Dataset

<https://www.kaggle.com/what0919/intrusion-detection>

Data basically represents the packet data for a time duration of 2 seconds.

1-9 Columns: basic features of packet (type 1)

10-22 columns: employ the content features (type 2)

23-31 columns: employ the traffic features with 2 seconds of time window (type 4)

32-41 columns: employ the host based features

31

31

Let's go to the Coding Demo...

32

To be continued in the next session.....